



Education

Sagar Institute of Research and Technology, Bhopal, MP, India B.Tech. RGPV Artificial Intelligence and Machine Learning CGPA: 8.03	June 2026
St. Joseph's Convent S.S. School, Sagar, MP, India 12th PCM CBSE 84.2%	June 2022
St. Joseph's Convent S.S. School, Sagar, MP, India 10th CBSE 76.5%	June 2020

Technologies

Languages: Python, SQL, C++	Databases: MongoDB, MySQL, PostgreSQL
AI/GenAI: LLMs, RAG, AI Agents, MCP, Prompt Engineering	Infrastructure: Docker, RunPod, GCP, Azure
Frameworks: FastAPI, Flask, LangChain, LiveKit	AI Infrastructure: PyTorch, vLLM, llama.cpp, ollama

Projects

AI Voicebot Platform – LiveKit

Technologies: Python, LiveKit, LLMs, STT/TTS, SIP, FastAPI, Docker, RAG, FAISS, Pinecone, ChromaDB, pgvector

- **AI Model & Voicebot Architecture:** Built a configurable **real-time AI voicebot** architecture supporting multiple STT, TTS, and LLM providers/models, with telephony/SIP integration and flexible deployment across multiple client use cases.
- **Business Integrations & AI Intelligence:** Integrated Shopify, Google Calendar, email, SMS, CRM systems, and third-party APIs, while implementing post-call pipelines for transcription, speaker diarization, summarization, key-point and entity extraction, conversational insights, and RAG-based analytics using **FAISS, Pinecone, ChromaDB, and pgvector**.
- **Production Scale & Operations:** Managed deployment, monitoring, maintenance, troubleshooting, and performance optimization across multiple client implementations, with the voicebot handling up to **500 calls daily** and resolving L1-level customer queries to reduce the operational workload of client teams.

Self-Hosted LLM & GPU Inference Platform

Technologies: Qwen, Gemma, Ollama, llama.cpp, vLLM, RunPod, NVIDIA A100, PyTorch, CUDA, FP8, LLM Inference APIs

- **Self-Hosted LLM Infrastructure:** Designed and deployed a local LLM inference architecture on **RunPod**, hosting open-source models on GPU infrastructure and exposing them through API-based inference services.
- **Inference & Model Evaluation:** Evaluated **Ollama, llama.cpp, and vLLM** for latency, concurrency, throughput, and resource utilization, with **vLLM** delivering the strongest results; light-weight benchmarked models including **Qwen3.8-27B-FP8, Qwen3, Gemma, Llama, and DeepSeek** for reasoning quality and deployment efficiency.
- **LLM Application Architecture:** Built a complete LLM-powered chatbot infrastructure with **MCP-based tool calling**, back-end integrations, and data retrieval, and optimized the architecture for **call analytics/auditing workloads**, with a deeper focus on local LLM deployment, inference optimization, and production-oriented GPU utilization.

Experience

AI Developer – Triotech Bizserve Pvt. Ltd.

Apr 2026 – Present

Technologies: Python, FastAPI, LLMs, RAG, LiveKit, MongoDB, SQL, Docker, RunPod, PyTorch, CUDA, vLLM

Projects & Responsibilities:

- **Applications Worked On:** Worked across production AI applications including real-time AI voicebots, call analytics, RAG-based chatbots, self-hosted LLM infrastructure, and AI-powered image analysis and automation workflows.
- **Voicebot Development & Optimization:** Developed and optimized a configurable real-time AI voicebot platform using LiveKit, LLMs, STT/TTS, and telephony/SIP, integrating CRM systems, Shopify, Google Calendar, email, SMS, and third-party APIs, with deployments handling up to **500 calls daily** and L1-level customer queries.
- **Technical Ownership & Client Engagement:** Served as the **technical anchor** in client meetings, bridging non-technical stakeholders and the development team, translating business requirements into technical solutions, coordinating implementations, and supporting client-facing technical discussions.
- **LLM & AI Engineering:** Worked extensively with LLMs, RAG, AI agents, MCP, prompt engineering, STT/TTS, and self-hosted inference, evaluating models and infrastructure for reasoning, tool calling, accuracy, latency, concurrency, and production workloads.
- **Teamwork & Leadership:** Collaborated closely with product, onboarding, development, and client teams, taking ownership of assigned deliverables, supporting teammates, resolving production issues, and driving tasks through development, deployment, and optimization with a strong focus on reliability and client requirements.

Achievements & Leadership

- **GDG Campus Ambassador:** Organized developer events and mentored students in AI and software projects.
- **Academic Hackathons:** Secured **3rd Place** in the 1 Billion Row Analysis competition at **IIT-BHU** and **5th Place** in a national-level hackathon.
- **Academic Research & IP:** Developed an ML-based **Cardiovascular Disease Prediction** solution with registered copyright (**L-157174/2024**) and presented research at conferences including **IIT Mandi**.
- **Certification:** **Python for Data Science - NPTEL** (National Programme on Technology Enhanced Learning).